

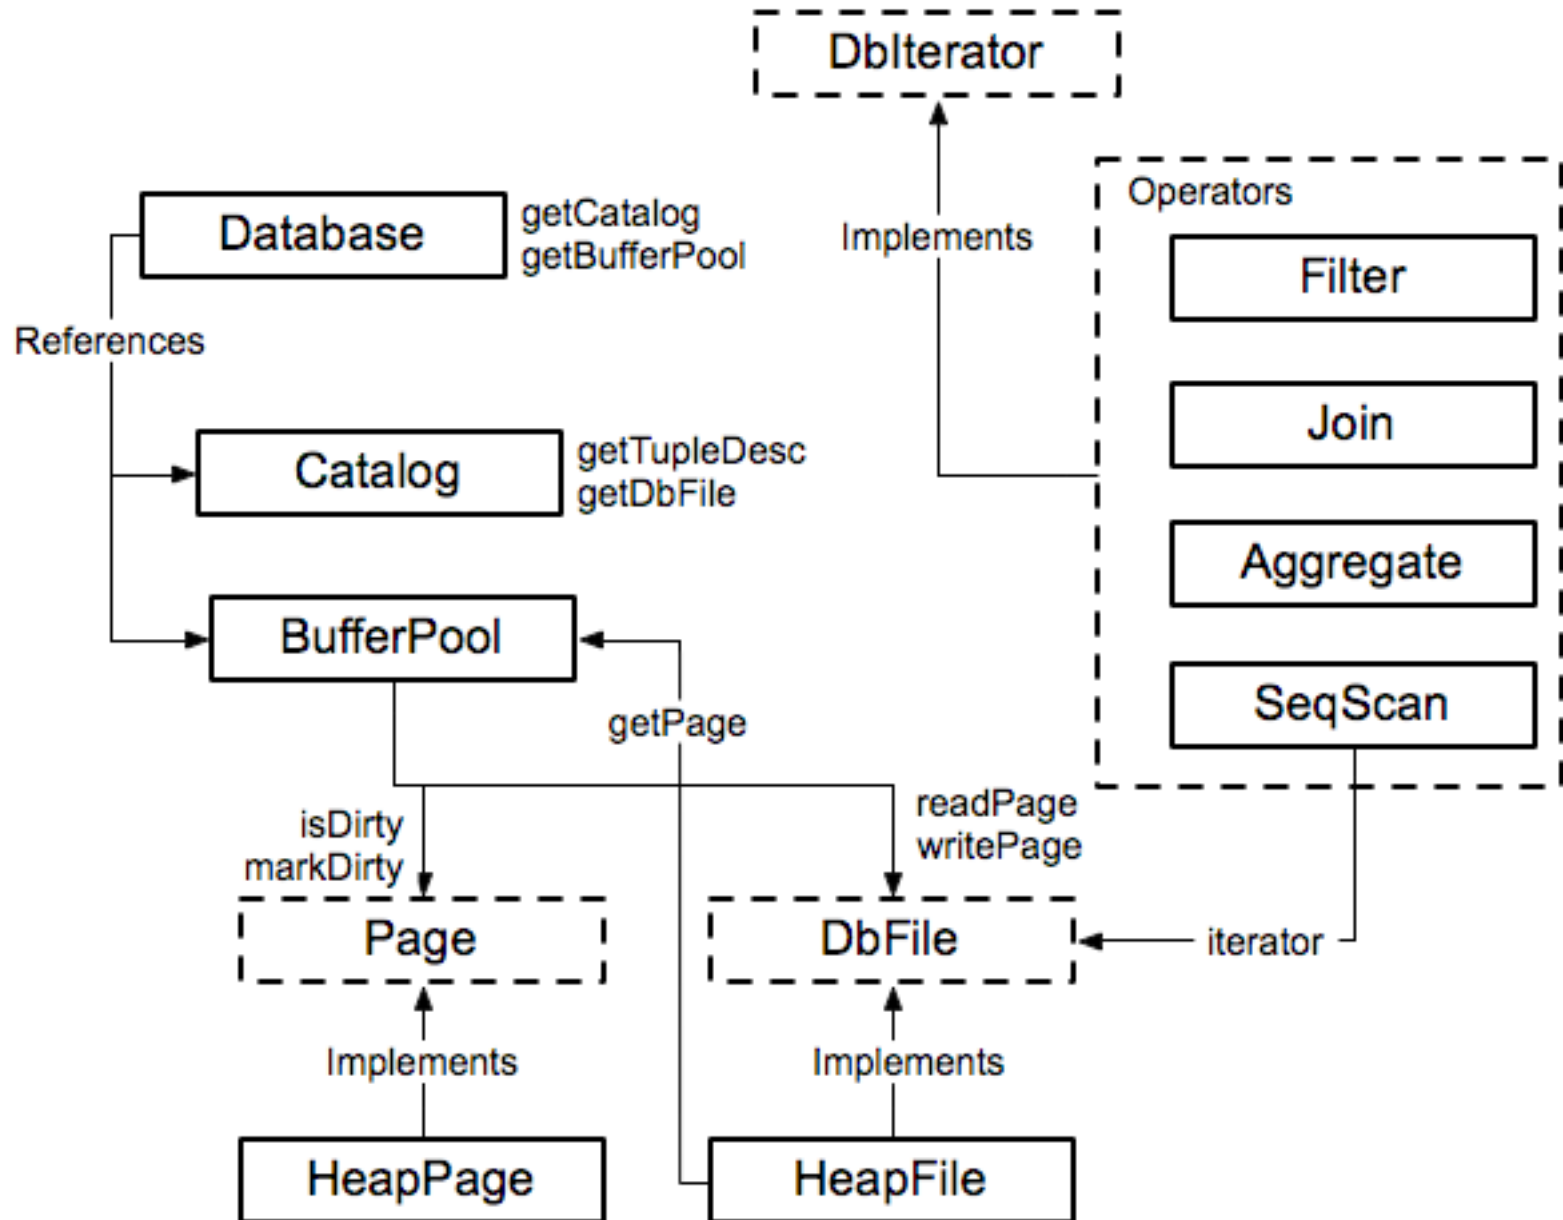
# SimpleDB Overview

9/18/2008

# What is SimpleDB?

- A basic database system
- What it has
  - Heapfiles
  - Basic Operators (Scan, Filter, JOIN, Aggregate)
  - Buffer Pool
  - Transactions
  - SQL Front-end
- Things it doesn't have
  - Query optimizer
  - Fancy relational operators (UNION, etc)
  - Recovery
  - Indices

# Module Diagram



# Catalog

- Catalog stores a list of available tables, TupleDesc
  - void addTable(DbFile d, TupleDesc d)
  - DbFile getTable(int tableid)
  - TupleDesc getTupleDesc(int tableid)
- Not persisted to disk

# DbIterator.java

- Iterator class implemented by all operators
  - open()
  - close()
  - getTupleDesc()
  - hasNext()
  - next()
  - rewind()
- Iterator model: chain iterators together

*// construct a 3-column table schema*

```
Type types[] = new Type[]{ Type.INT_TYPE, Type.INT_TYPE, Type.INT_TYPE };
```

```
String names[] = new String[]{ "field0", "field1", "field2" };
```

```
TupleDesc descriptor = new TupleDesc(types, names);
```

*// create the table, associate it with some\_data\_file.dat*

*// and tell the catalog about the schema of this table.*

```
HeapFile table1 = new HeapFile(new File("some_data_file.dat"), descriptor);
```

```
Database.getCatalog().addTable(table1);
```

*// construct the query: we use a simple SeqScan, which spoonfeeds*

*// tuples via its iterator.*

```
TransactionId tid = new TransactionId();
```

```
SeqScan f = new SeqScan(tid, table1.id());
```

*// and run it*

```
f.open();
```

```
while (f.hasNext()) {
```

```
    Tuple tup = f.next();
```

```
    System.out.println(tup);
```

```
}
```

```
f.close();
```

```
Database.getBufferPool().transactionComplete();
```

# HeapFile.java

- An array of HeapPages on disk
- Javadoc is your friend!
- Implement everything except addTuple and removeTuple

# HeapPage.java

- Format
  - Header is a bitmap
  - Page contents are an array of fixed-length Tuples
- Full page size = `BufferPool.PAGE_SIZE`
- Number of bits in Header = number of Tuples
- Header size + size of tuples = `BufferPool.PAGE_SIZE`



# HeapFileEncoder.java

- Because you haven't implemented insertTuple, you have no way to create data files
- HeapFileEncoder converts CSV files to HeapFiles
- Usage:
  - `java -jar dist/simpliedb.jar convert csv-file.txt numFields`
- Produces a file `csv-file.dat`, that can be passed to HeapFile constructor.

# BufferPool.java

- Manages cache of pages
  - Evicts pages when cache is full [not lab 1]
- All page accesses should use getPage
  - Even from inside DbFile!

You will eventually implement

- locking for transactions
- Flushing of pages for recovery

# Compiling, Testing, and Running

- Compilation done through the ant tool
  - Works a lot like make
- Two kinds of tests:
  - Unit tests
  - System Tests
- Demo on debugging using unit tests.