

6.830 Problem Set 3 (2014)

Assigned: Monday, October 27, 2014

Due: December 1, 2014, 11:59 PM

Submit to Stellar: <https://stellar.mit.edu/S/course/6/fa14/6.830/homework/>)

Revision history:

Nov 20, 2014 – Updated Problem 3 to replace `WRITE A, x2 + x1` with `WRITE A, x2 + y2` in T2.

Nov 26, 2014 – Updated Problem 10 so that worker logs aren't impossible.

The purpose of this problem set is to give you some practice with concepts related to recovery, parallel query processing, two-phase commit, and other papers we read during the second half of the course.

1 Questions

1. [5 points]: Below, we provide 3 workloads. Each one contains several concurrent transactions (consisting of `READ` and `WRITE` statements on objects). For each workload, several possible execution interleavings are given. Your job is to indicate whether, for each of the interleavings:

- The interleaving has an equivalent serial ordering. If so, indicate what the serial ordering is.
- Whether the interleaving would be valid using lock-based concurrency control. Assume that locks, when needed, are acquired (with the appropriate lock mode) on an object just before the statement that reads or writes the object and that all locks are released during the `COMMIT` statement (and no sooner.) If the interleaving is not valid, indicate whether or not it would simply never occur or would result in deadlock, and the time when the deadlock would occur.
- Whether the ordering would be valid using optimistic concurrency control. If not, indicate which transaction will be aborted. Assume the use of the Parallel Validation scheme described in Section 5 of the Optimistic Concurrency Control paper by Kung and Robinson, and that the validation and the write phases of optimistic concurrency control happen during the `COMMIT` statement (and no sooner.)

Assume in all cases that written values can depend on previously read values. (The workloads are shown on the next page.)

Workload 1

<u>Transaction 1</u>	<u>Transaction 2</u>
READ A	READ A
READ B	WRITE B
WRITE C	WRITE A

Interleaving 1:

```

1 T1:  READ A
2 T2:  READ A
3 T1:  READ B
4 T1:  WRITE C
5 T2:  WRITE B
6 T1:  COMMIT
7 T2:  WRITE A
8 T2:  COMMIT

```

Interleaving 2:

```

1 T1:  READ A
2 T1:  READ B
3 T2:  READ A
4 T2:  WRITE B
5 T2:  WRITE A
6 T2:  COMMIT
7 T1:  WRITE C
7 T1:  COMMIT

```

Workload 2

<u>Transaction 1</u>	<u>Transaction 2</u>	<u>Transaction 3</u>
READ A	READ A	READ A
WRITE A	WRITE B	WRITE A

Interleaving 3:

```

1 T1:  READ A
2 T2:  READ A
3 T3:  READ A
4 T2:  WRITE B
5 T2:  COMMIT
6 T3:  WRITE A
7 T3:  COMMIT
8 T1:  WRITE A
9 T1:  COMMIT

```

Workload 3

<u>Transaction 1</u>	<u>Transaction 2</u>
WRITE A	WRITE B
READ B	READ C

Interleaving 4:

```

1 T1:  Write A
2 T2:  Write B
3 T1:  Read B
4 T2:  Read C
5 T1:  Commit
6 T2:  Commit

```

Interleaving 5:

```

1 T1:  Write A
2 T2:  Write B
3 T2:  Read C
4 T2:  Commit
5 T1:  Read B
6 T1:  Commit

```

2. [3 points]: If your DBMS never STOLE pages (as discussed in class and in the paper by Franklin), how would that affect the design of the database recovery manager?

3. [3 points]: Suppose you are told that the following transactions are run concurrently on a (locking-based, degree 3 consistency) database system that has just been restarted and is fully recovered. Suppose the system crashes while executing the statement marked by an “***” in Transaction 1. Suppose that Transaction 2 has committed, and the state of Transactions 3 and 4 are unknown (e.g., they may or may not have committed.) Assume that each object (e.g., X, Y, etc.) occupies exactly one page of memory.

<u>Trans 1:</u>	<u>Trans 2:</u>	<u>Trans 3:</u>	<u>Trans 4:</u>
x1 = READ X	WRITE Y, 0	z3 = READ X	a4 = READ A
WRITE X, x1 + 1	WRITE B, 0	a3 = READ A	z4 = READ Z
*** y1 = READ Y	x2 = READ X	WRITE A, a3 + 10	WRITE B, (a4 - z4)
WRITE Y, y1 + x	WRITE Z, x2	z3 = READ Z	
	y2 = READ Y	WRITE Z, z3 - 10	
	WRITE A, x2 + y2		

- Show an equivalent serial order that could have resulted from these statements, given what you know about what statement was executing when the system crashed. In addition, show an interleaving of the statements from these transactions that is equivalent to your serial order; make sure this serial order could result from a locking-based concurrency control protocol (again assuming that locks are acquired immediately before an item is accessed and released just before the commit statement.)
- Show all of the records that should be in the log at the time of the crash (given your serial order), assuming that there have been no checkpoints and that you are using an ARIES-style logging and recovery protocol. Your records should include all of the relevant fields described in Section 4.1 of the ARIES paper. Also show the status of the transaction table (as described in Section 4.3 of the ARIES paper) after the analysis phase of recovery has run.
- Suppose you have 2 pages of memory, and are using a STEAL/NO-FORCE buffer management policy as in ARIES. Given the interleaving you showed above, for each of the 5 pages used in these transactions, show one possible assignment of LSN values for those pages as they are on disk before recovery runs. You should use the value “?” if the LSN is unchanged from the prior state of the page before these transactions began. Finally, indicate which pages will be modified during the UNDO pass, and which will be modified during the REDO pass.

2 Parallel Query Processing

The standard way to execute a query in a shared-nothing parallel database is to horizontally partition all tables across all nodes in the database. Under such a setting, distributed joins can be computed by repartitioning tables over the join attributes.

```
SELECT *
FROM R, S
WHERE R.a > v1 AND S.b > v2 AND R.c = S.d
```

You are given the following

- Both tables are 6,000 MB,
- The disk can read at 50 MB/sec (for the purposes of this problem, you may ignore differences between sequential and random I/O),
- The network on each node can transmit data at 40 MB/sec, regardless of the number or rate at which other nodes are simultaneously transmitting,
- A computer cannot send over the network and read or write from its disk at the same time,

- The selectivity of both selection predicates is 0.05,
- Each tuple in R joins with exactly one tuple in S,
- Each machine in your distributed database has 300 MB of memory.

Suppose both tables are stored on a single node.

4. [3 points]: Describe the best query plan for executing this query on that node.

5. [2 points]: Ignoring CPU costs, estimate the time to answer this query on a single node.

Now, suppose that the tables are hash-partitioned on R.a and S.b across a 4 node distributed database.

6. [3 points]: Describe the best distributed query plan for executing this query.

7. [2 points]: Ignoring CPU costs, estimate the time to answer this query on the distributed database.

Now, suppose that the tables are hash-partitioned on R.c and S.d across a 4 node distributed database.

8. [3 points]: Describe the best distributed query plan for executing this query.

9. [2 points]: Ignoring CPU costs, estimate the time to answer this query on the distributed database.

3 Two Phase Commit

Recall that the two-phase commit protocol is used to process transactions over a distributed database on several nodes. In this set of questions, suppose you are running the following three transactions over five values, A, B, C, D, and E, which initially have value 0. Assume the use of the basic (e.g., not presumed commit or presumed abort version) of the protocol described in the paper “Transaction Management in the R* Distributed Database Management Systems” by C. Mohan et al.

```
T1:
WA(1)
WB(2)
WE(3)
COMMIT
```

```
T2:
WB(4)
WC(5)
WE(6)
COMMIT
```

```
T3:
WC(7)
WD(8)
WE(9)
COMMIT
```

Here, the notation WA(1) means “write value 1 to A”. You run these three transactions concurrently on three workers, W1, W2 and W3, as well as a coordinator, C. Data items A and C are on W1, B and D are on W2, E is on W3. The system crashes during the execution, and you know that each transaction has begun but do not know how far into its execution it has proceeded. Before recovery begins, you observe the a small piece of the logs on one or more of the workers and the coordinator. Based on just the information in each set of partial logs shown below, indicate whether each transaction has already or definitely will commit or abort, or whether its outcome is unknown.

10. [3 points]:

W1 log:

```
...
PREPARE T1 -- VOTE YES
...
```

W2 log:

```
...
T3 UP D
...
COMMIT T3
...
```

W3 log:

```
...
T2 UP E
...
PREPARE T3 -- VOTE YES
COMMIT T2
...
```

There may be other log entries in place of the ellipses (...) shown here. The notation “T2 UP E” means that T2 updated value E. Indicate whether each of T1, T2, or T3 will commit, abort, or its outcome is unknown.

11. [2 points]:

C Log:

```
...  
COMMIT T3  
...
```

W2 Log:

```
...  
PREPARE T1 -- VOTE NO  
...  
ACK T2  
...  
ACK T3  
....
```

Indicate whether each of T1, T2, or T3 will commit, abort, or its outcome is unknown.